

# Assignment 2

## CSE 130-02: Principles of Computer System Design, Fall 2019

Due: Wednesday, November 20 at 9:00PM

### Goals

The goals for Assignment 2 are to modify the HTTP server that you already implemented to have two additional features: multi-threading and logging. Multi-threading that your server must be able to handle multiple requests simultaneously, each in its own thread. Logging means that your server must write out a record of each request, including both header information and data (dumped as hex). You'll need to use synchronization techniques to service multiple requests at once, and to ensure that entries in the log aren't intermixed from multiple threads.

As usual, you must have a design document and writeup along with your `README.md` in your `git` repository. Your code must build `httpserver` using `make`.

### Programming assignment: multi-threaded HTTP server with logging

#### Design document

Before writing code for this assignment, as with every other assignment, you must write up a design document. Your design document must be called `DESIGN.pdf`, and must be in PDF. You can easily convert other document formats, including plain text, to PDF. Scanned-in design documents are fine, as long as they're legible and they reflect the code you *actually* wrote.

Your design should describe the design of your code in enough detail that a knowledgeable programmer could duplicate your work. This includes descriptions of the data structures you use, non-trivial algorithms and formulas, and a description of each function with its purpose, inputs, outputs, and assumptions it makes about inputs or outputs.

Write your design document *before* you start writing code. It'll make writing code a lot easier. Also, if you want help with your code, the first thing we're going to ask for is your design document. We're happy to help you with the design, but we can't debug code without a design any more than you can.

You **must** commit your design document *before* you commit the code specified by the design document. You're welcome to do the design in pieces (*e.g.*, multi-threading first, then the design for logging), as long as you don't write code for the parts that aren't well-specified. We **expect** you to commit multiple versions of the design document; your commit should specify *why* you changed the design document if you do this (*e.g.*, "original approach had flaw X", "detailed design for module Y", etc.). **If you commit code before it's designed, or you commit your design a few minutes before the working code that the design describes, you will lose points.** We want you to get in the habit of designing components before you build them.

Since a lot of the system in Assignment 2 is similar to Assignment 1, we expect you're going to "copy" a good part of your design from your Assignment 1 design. This is fine, as long as it's *your* Assignment 1 you're copying from. This will let you focus on the new stuff in Assignment 2.

For multithreading, getting the design of the program right is vital. We expect the design document to contain discussions of *why* your system is thread-safe. Which variables are shared between threads? When are they modified? Where are the critical regions? These are some of the things we want to see. Similarly, you should explain why your design for logging each request contiguously actually works.

#### Program functionality

Your code may be either C or C++, but all source files must have a `.cpp` suffix and be compiled by `clang++` version 7. As before, you may not use standard libraries for HTTP, nor any `FILE *` or `iostream` calls except for printing to the screen (*e.g.*, error messages). You may use standard networking (and file system) system calls.

Details on the HTTP functionality your server must support are available in Assignment 1. Assignment 2 has two minor changes, in addition to multithreading and logging. First, PUT headers now *must* include a `Content-Length`:

header; a PUT request that is missing such a header is a bad request. Second, resource names may optionally be preceded by a /. The name (excluding the slash, if present) must still be 27 characters long consisting only of the 64 allowed characters from Assignment 1. This means that `abcdefghij0123456789abcdefgh` and `/abcdefghij0123456789abcdefgh` are both legal, and both refer to the same file, whose name is `abcdefghij0123456789abcdefgh`.

We expect that you'll build on your server code from Assignment 1 for Assignment 2. Remember, however, that your code for this assignment must be developed in `asgn2`, so copy it there before you start, and make sure you add it to your repository using `git add`.

## Multi-threading

Your previous Web server could only handle a single request at a time, limiting throughput. Your first goal for Assignment 2 is to use *multi-threading* to improve throughput. This will be done by having each request processed in its own thread. The typical way to do this is to have a “pool” of “worker” threads available for use. The server creates  $N$  threads when it starts;  $N = 4$  by default, but the argument `-N nthreads` tells the server to use  $N = nthreads$ . For example, `-N 6` would tell the server to use 6 worker threads.

Each thread waits until there is work to do, does the work, and then goes back to waiting. Worker threads may not “busy wait” or “spin lock”; they must actually sleep by waiting on a lock, condition variable, or semaphore. Each worker thread does what your server did for Assignment 1 to handle client requests, so you can likely reuse much of the code for it. However, you'll have to add code for the dispatcher to pass the necessary information to the worker thread. Worker threads should be independent of one another; if they ever need to access shared resources, they must use synchronization mechanisms for correctness. Each thread may allocate up to 16 KiB of buffer space for its own use; this includes space for send/receive buffers as well as log entry buffers (see below).

A single “dispatch” thread listens to the connection and, when a connection is made, alerts (using a synchronization method such as a semaphore or condition variable) one thread in the pool to handle the connection. Once it has done this, it assumes that the worker thread will handle the connection itself, and goes back to listening for the next connection. The dispatcher thread is a small loop (likely in a single function) that contacts one of the threads in the pool when a request needs to be handled. If there are no threads available, it must wait until one is available to hand off the request. Here, again, a synchronization mechanism must be used.

You will be using the POSIX threads library (`libpthreads`) to implement multithreading. There are a lot of calls in this library, but you only need a few for this assignment. You'll need:

- `pthread_create (...)`
- Condition variables, mutexes and/or semaphores. You may use any or all, as you see fit. See `pthread_cond_wait`, `pthread_cond_signal`, `pthread_cond_init`, `pthread_mutex_init`, `pthread_mutex_lock`, `pthread_mutex_unlock`, `sem_init`, `sem_overview`, `sem_wait`, and `sem_post` for details.

Since your server will never exit, you don't need `pthread_exit`, and you likely won't need `pthread_join` since your thread pool never gets smaller. You may not use any synchronization primitives other than (basic) semaphores, locks, and condition variables. You are, of course, welcome to write code for your own higher-level primitives from locks and semaphores if you want.

There are several `pthread` tutorials available on the internet, including this one (has multiple pages—see the next page button at the bottom), and this one. We will also cover some basic `pthreads` techniques in section.

If you prefer, you may use the C++ STL threading library, but *only* the classes in `<thread>`, `<mutex>`, and `<condition_variable>`.

## Logging requests

If the `-l log_file` option is provided to your program, it must log each request to the file `log_file` (obviously, `log_file` is whatever name you provide on the command line). If the `-l` option isn't provided, no logging is done.

A log record looks like this:

```
PUT abcdefghij0123456789abcdefg length 36
00000000 65 68 6c 6c 3b 6f 68 20 6c 65 6f 6c 61 20 61 67 6e 69 20 3b
00000020 65 68 6c 6c 20 6f 54 28 65 68 43 20 72 61 29 73
=====
```

The first line contains the operation (PUT or GET), along with the file name and the length of the data. You'll know this before you receive all the data from the `Content-Length` header, or before you send the data from the information you're going to print in the response header. The first line is followed by lines that start with a zero-padded byte count, followed by up to 20 bytes printed as hex numbers. Your program can use `+sprintf()` to format each line into a buffer, and then write out the buffer to the file as it gets full. Lines in the log end in `\n`, *not* `\r\n`, so (for example) the the second line contains  $8 + 20 \times 3 + 1 = 69$  characters. The log entry is terminated by `=====``\n`.

If the server returns an error response code, the log record looks like this:

```
FAIL: GET abcd HTTP/1.1 --- response 400\n
```

The record starts with `FAIL:`, followed by the first line of the header (without the `\r\n` at the end), followed by `" --- response XXX\n"`, where XXX is the response code (without the explanation) that your server returned.

Each request must be logged *contiguously*. This can be tricky, since there are multiple threads writing to the log file at the same time, which might cause the log entries to get interleaved. To avoid interleaving requests in the log file, a thread should “reserve” the space it's going to use, and use `pwrite(2)` to write data to a specific location in the file. Your program will need to use synchronization between threads to reserve the space, but should need no synchronization to actually write the log information, since no other thread is writing to that location. Remember, if you know the length of the data being sent or received, you can calculate exactly how much space you need for the request log entry.

## README and Writeup

Your repository must also include a README file (`README.md`) and writeup (`WRITEUP.pdf`). The README may be in either plain text or have Markdown annotations for things like bold, italics, and section headers. **The file must always be called `README.md`**; plain text will look “normal” if considered as a Markdown document. You can find more information about Markdown at <https://www.markdownguide.org>.

The `README.md` file should be short, and contain any instructions necessary for running your code. You should also list limitations or issues in `README.md`, telling a user if there are any known issues with your code.

Your `WRITEUP.pdf` is where you'll describe the testing you did on your program and answer any short questions the assignment might ask. The testing can be unit testing (testing of individual functions or smaller pieces of the program) or whole-system testing, which involves running your code in particular scenarios.

For Assignment 2, please answer the following question:

- Using either your HTTP client or `curl` and your *original* HTTP server from Assignment 1, do the following:
  - Place four different large files on the server. This can be done simply by copying the files to the server's directory, and then running the server. The files should be around 4 MiB long.
  - Start `httpserver`.

- Start four *separate* instances of the client *at the same time*, one GETting each of the files and measure (using `time(1)`) how long it takes to get the files. Perhaps the best way to do this is to write a simple shell script (command file) that starts four copies of the client program in the background, by using `&` at the end.
- Repeat the same experiment after you implement multi-threading. Is there any difference in performance?
- What is likely to be the bottleneck in your system? How much concurrency is available in various parts, such as dispatch, worker, logging? Can you increase concurrency in any of these areas and, if so, how?

## Submitting your assignment

All of your files for Assignment 2 must be in the `asgn2` directory in your `git` repository. When you push your repository to GITLAB@UCSC, the server will run a program to check the following:

- There are no “bad” files in the `asgn2` directory (*i.e.*, object files).
- Your assignment builds in `asgn2` using `make` to produce `httpserver`.
- All required files (`DESIGN.pdf`, `README.md`, `WRITEUP.pdf`) are present in `asgn2`.

If the repository meets these minimum requirements for Assignment 2, there will be a green check next to your commit ID in the GITLAB@UCSC Web GUI. If it doesn’t, there will be a red X. **It’s OK to commit and push a repository that doesn’t meet minimum requirements for grading.** However, we will only *grade* a commit that meets these minimum requirements.

Note that the *minimum* requirements say nothing about correct functionality—the green check only means that the system successfully ran `make` and that all of the required documents were present, with the correct names.

**You must submit the commit ID you want us to grade via Google Form, linked to the assignment page on Canvas. This must be done before the assignment deadline.**

## Hints

- Start early on the design. This program builds on Assignment 1. If you didn’t get Assignment 1 to work, please see the course staff ASAP for help getting it to work.
- Reuse your code from Assignment 1. No need to cite this; we expect you to do so.
- Go to section for additional help with the program. This is especially the case if you don’t understand something in this assignment!
- You’ll need to use (at least) the system calls from Assignment 1, as well as `pthread_create` and some form of mutual exclusion (semaphores and/or mutexes). You’ll also need `pwrite(3)`, which is like `write`, but takes a file offset.
- Test multi-threading and logging separately before trying them together.
- Aggressively check for and report errors. Transfers may be aborted on errors. However, the server doesn’t *exit* on an error; it just goes on to the next transfer.
- Use `getopt(3)` to parse options from the command line. Read the man pages and see examples on how it’s used. Ask the course staff if you have difficulty using it after reading this material.
- A sample command line for the server might look like this:  
`./httpserver -N 8 -l my_log.txt localhost 8888`  
This would tell the server to start 8 threads, and to write log records to `my_log.txt`. The server would listen for connections on `localhost`, port 8888.

## Grading

As with all of the assignments in this class, we will be grading you on *all* of the material you turn in, with the *approximate* distribution of points as follows: design document (35%); coding practices (15%); functionality (40%); writeup (10%).

**If you submit a commit ID without a green checkmark next to it or modify `.gitlab-ci.yml` in any way, your maximum grade is 5%. Make sure you submit a commit ID with a green checkmark.**

## Extra Credit

We will award 15 points of extra credit to the 5 fastest servers in the class, and 10 points to the next 15 fastest servers. Please note that your Web server must pass all of the tests in the GitLab test framework to be eligible for extra credit, and you must have at least a 70 on the assignment to be eligible. The extra credit may raise your grade above 100 points.